

# Salary Prediction Project Report (Machine Learning 1/2025)

## 1. Objective

To develop an accurate classification model predicting high-income individuals using census demographic and employment data, comparing multiple machine learning approaches to identify the optimal solution for income level classification and deployment.

## 2. Data Understanding

The dataset consists of **20,900 census records** with **18 original features**, used to predict whether an individual earns a high salary based on demographic and employment-related factors

- Sample Size: 16,720 training records, 4,180 test records, 6,967 live prediction records
- Target Variable: Binary label (1.0 = high income, 0.0 = low income)
- Feature Types: Numerical (age-group, education-num, capitalgain, hoursperweek) and categorical (workclass, education, occupation, race, sex)
- Data Quality: Previously addressed missing values (~5% in workclass/occupation) and applied standardization/encoding

## 3. Preparation

Data Processing Pipeline:

1. Data Splitting: Separate training to 70% and test sets to 30% for model validation
2. Handle Missing Categorical Data: Used **SimpleImputer(strategy='most\_frequent')** to fill missing values.
3. Handle Missing Numerical Data: Applied **SimpleImputer(strategy='mean')** on numerical columns.
4. Encode Ordinal Features: Applied **OrdinalEncoder** for education, ranking from lowest (preschool) to highest (doctorate).
5. Encode Nominal Features: Used **OneHotEncoder** for workclass, occupation, relationship, and sex
6. Feature Scaling (Standardization): Used **StandardScaler** on numerical variables like age-group, education-num, capitalgain, capitalloss, and hoursperweek.

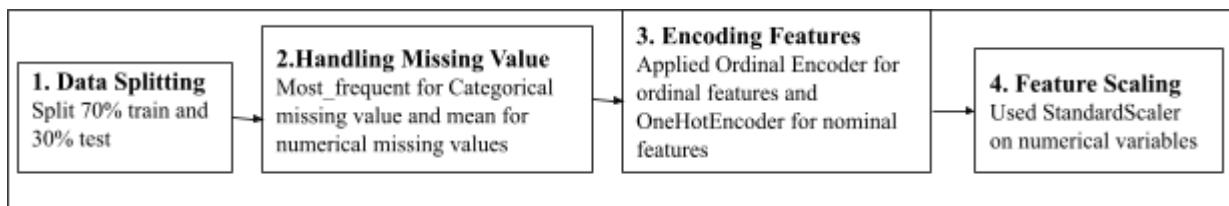


Figure 1: diagram showing data preparation process.

## 4. Modeling Approach

1. **Model selection rationale:** The **Artificial Neural Network (ANN)** was chosen for its ability to capture **non-linear relationships** between demographic, educational, and occupational features. It performs better than simple linear models for complex socio-economic data.

2. **Model configuration/hyperparameters:** The model was built using `sklearn.neural_network.MLPClassifier` with these main settings: hidden layers = (20, 10), activation = ‘logistic’, solver = ‘sgd’, learning\_rate\_init = 0.1, batch\_size = 32, max\_iter = 1000, and random\_state = 0. All categorical features were one-hot encoded, and numerical features were scaled using **MinMaxScaler**.
3. **Experimental Setup:** The dataset was split into 70% training and 30% testing using `train_test_split`. Input features ( $X_{train}$ ) were derived from cleaned and scaled attributes; the target variable ( $y_{train}$ ) was binary(0.0 = low salary, 1.0 = high salary). Model fitting used the training set only, while evaluation employed the test set for unbiased performance estimation. Although a convergence warning appeared after 1000 iterations, the model reached stable loss and strong generalization.
4. **Special Techniques:** We applied feature scaling for faster convergence, one-hot encoding for categorical variables, and manual tuning of learning rate and hidden layers to balance model complexity and training time.

## 5. Evaluation and Results

	<b>0.0</b>	<b>1.0</b>	<b>accuracy</b>	<b>macro avg</b>	<b>weighted avg</b>
<b>precision</b>	0.871718	0.804015	0.842436	0.837866	0.843246
<b>recall</b>	0.853722	0.826886	0.842436	0.840304	0.842436
<b>f1-score</b>	0.862626	0.815290	0.842436	0.838958	0.842719
<b>support</b>	8477.000000	6152.000000	0.842436	14629.000000	14629.000000

Figure 2: Table report performance metrics.

Model	Accuracy	Precision	Recall	F1-Score	support
ANN	0.842436	0.804015	0.82688	0.815290	6152.000000
RF	0.814543	0.759485	0.812093	0.784908	2613.000000
KNN	0.790942	0.734510	0.734510	0.756727	2613.000000

Figure 3: Table comparison of multiple models.

## 6. Discussion and Analysis

The **ANN model** achieved the best F1-score and accuracy, effectively classifying salary levels. It detected high-salary cases well but with slightly lower precision. Minor overfitting may occur due to long training, yet test results remain stable. Future work should apply **cross-validation**, **regularization**, and **learning rate tuning** to improve stability and interpretability.